

Write a function `ICKFib()` that takes an integer n and returns the n th Fibonacci number.

Let's say our Fibonacci series is 0-indexed and starts with 0. So:

```
ICKFib(0); // => 0
ICKFib(1); // => 1
ICKFib(2); // => 1
ICKFib(3); // => 2
ICKFib(4); // => 3
...
```

Objective-C ▼

Gotchas

Our solution runs in n time.

There's a clever, more mathy solution that runs in $O(\lg n)$ time, but we'll leave that one as a bonus.

If you wrote a recursive function, think carefully about what it does. It might do repeat work, like computing `ICKFib(2)` multiple times!

We can do this in $O(1)$ space. If you wrote a recursive function, there might be a hidden space cost in the call stack!

Breakdown

The n th Fibonacci number is defined in terms of the two *previous* Fibonacci numbers, so this seems to lend itself to recursion.

```
ICKFib(n) = ICKFib(n - 1) + ICKFib(n - 2);
```

Objective-C ▼

Can you write up a recursive solution?

As with any recursive function, we just need a base case and a recursive case:

1. **Base case:** n is 0 or 1. Return n .
2. **Recursive case:** Return $\text{ICKFib}(n - 1) + \text{ICKFib}(n - 2)$.

```
NSInteger ICKFib(NSUInteger n) {  
  
    if (n == 0 || n == 1) {  
        return n;  
    }  
  
    return ICKFib(n - 1) + ICKFib(n - 2);  
}
```

Objective-C ▼


Okay, this'll work! What's our time complexity?

It's not super obvious. We might guess n , but that's not quite right. Can you see why?

Each call to `ICKFib()` makes *two more calls*. Let's look at a specific example. Let's say $n = 5$. **If we call `ICKFib(5)`, how many calls do we make in total?**

Try drawing it out as a tree where each call has two child calls, unless it's a base case.

Here's what the tree looks like:



We can notice this is a binary tree, whose height is n , which means the total number of nodes is $O(2^n)$.

So our total runtime is $O(2^n)$. That's an "exponential time cost," since the n is *in an exponent*. Exponential costs are *terrible*. This is way worse than $O(n^2)$ or even $O(n^{100})$.

Our recurrence tree above essentially gets twice as big each time we add 1 to n . So as n gets really big, our runtime quickly spirals out of control.

The craziness of our time cost comes from the fact that we're doing so much repeat work. How can we avoid doing this repeat work?

We can memoize!

Let's wrap `ICKFib()` in a class with an instance variable where we store the answer for any n that we compute:

```

@interface ICKFibber : NSObject {
    @private
    NSMutableDictionary<NSNumber *, NSNumber *> *_memo;
}
- (NSUInteger)fib:(NSUInteger)n;
@end

@implementation ICKFibber

-(instancetype)init {
    if (self = [super init]) {
        _memo = [NSMutableDictionary new];
    }
    return self;
}

- (NSUInteger)fib:(NSUInteger)n {

    // base case: 0 or 1
    if (n == 0 || n == 1) {
        return n;
    }

    // see if we've already calculated this
    if (_memo[@(n)]) {
        return _memo[@(n)].unsignedIntegerValue;
    }

    NSUInteger result = [self fib:(n - 1)] + [self fib:(n - 2)];

    // memoize
    _memo[@(n)] = @(result);


    return result;
}

@end

```

What's our time cost now?

Our recurrence tree will look like this:



The computer will build up a call stack with `ICKFib(5)`, `ICKFib(4)`, `ICKFib(3)`, `ICKFib(2)`, `ICKFib(1)`. Then we'll start returning, and on the way back up our tree we'll be able to compute each node's 2nd call to `ICKFib()` in constant time by just looking in the memo. n time in total.

What about space? `_memo` takes up n space. Plus we're still building up a call stack that'll occupy n space. Can we avoid one or both of these space expenses?

Look again at that tree. Notice that to calculate `ICKFib(5)` we worked "down" to `ICKFib(4)`, `ICKFib(3)`, `ICKFib(2)`, etc.

What if instead we *started* with `ICKFib(0)` and `ICKFib(1)` and worked "up" to n ?

Solution

We use a bottom-up approach, starting with the 0th Fibonacci number and iteratively computing subsequent numbers until we get to n .

```
NSUInteger ICKFib(NSUInteger n) {  
  
    // edge cases:  
    if (n == 0 || n == 1) {  
        return n;  
    }  
  
    // we'll be building the fibonacci series from the bottom up  
    // so we'll need to track the previous 2 numbers at each step  
    NSUInteger prevPrev = 0; // 0th fibonacci  
    NSUInteger prev = 1;     // 1st fibonacci  
    NSUInteger current = 0;  // declare and initialize current  
  
    for (NSUInteger i = 1; i < n; i++) {  
  
        // Iteration 1: current = 2nd fibonacci  
        // Iteration 2: current = 3rd fibonacci  
        // Iteration 3: current = 4th fibonacci  
        // To get nth fibonacci ... do n-1 iterations.  
        current = prev + prevPrev;  
        prevPrev = prev;  
        prev = current;  
    }  
  
    return current;  
}
```

Complexity

$O(n)$ time and $O(1)$ space.

Bonus

- If you're good with matrix multiplication you can bring the time cost down even further, to $O(\lg(n))$. Can you figure out how?

What We Learned

This one's a good illustration of the tradeoff we sometimes have between code cleanliness and efficiency.

We could use a cute, recursive function to solve the problem. But that would cost $O(2^n)$ time as opposed to n time in our final bottom-up solution. Massive difference!

In general, whenever you have a recursive solution to a problem, think about what's *actually happening on the call stack*. An iterative solution might be more efficient.

Ready for more?

Check out our full course ➡

Want more coding interview help?

Check out **interviewcake.com** for more advice, guides, and practice questions.